



Assignment 1

The data has been loaded and stored in the drive and mounting drive.

```
[ ] #mounting drive for initializing dataframe
from google.colab import drive
drive.mount('/content/drive')

Mounted at /content/drive

[ ] #Loading dataset
data = pd.read_csv('/content/drive/MyDrive/Classroom/Adv. ML ML for DS/Assignment01/data.csv')

#printing head of dataset
print(data.head())

#printing description of dataset
print(data.describe())

#printing info of the dataset
data.info()
```

Head:

```
id diagnosis radius_mean texture_mean perimeter_mean area_mean \
0 842302 M 17.99 10.38 122.80 1001.0
1 842517 M 20.57 17.77 132.90 1326.0
2 84300903 M 19.69 21.25 130.00 1203.0
3 84348301 M 11.42 20.38 77.58 386.1
4 84358402 M 20.29 14.34 135.10 1297.0

smoothness_mean compactness_mean concavity_mean concave points_mean \
0 0.11840 0.27760 0.3001 0.14710
1 0.08474 0.07864 0.0869 0.07017
2 0.10960 0.15990 0.1974 0.12790
3 0.14250 0.28390 0.2414 0.10520
4 0.10030 0.13280 0.1980 0.10430

... texture_worst perimeter_worst area_worst smoothness_worst \
0 ... 17.33 184.60 2019.0 0.1622
1 ... 23.41 158.80 1956.0 0.1238
2 ... 25.53 152.50 1709.0 0.1444
3 ... 26.50 98.87 567.7 0.2098
4 ... 16.67 152.20 1575.0 0.1374

compactness_worst concavity_worst concave points_worst symmetry_worst \
0 0.6656 0.7119 0.2654 0.4601
1 0.1866 0.2416 0.1860 0.2750
2 0.4245 0.4504 0.2430 0.3613
3 0.8663 0.6869 0.2575 0.6638
4 0.2050 0.4000 0.1625 0.2364

fractal_dimension_worst Unnamed: 32
0 0.11890 NaN
1 0.08902 NaN
2 0.08758 NaN
3 0.17300 NaN
4 0.07678 NaN

[5 rows x 33 columns]
```

Description:

Info:

[5 rows x 33 columns]										[8 rows x 32 columns]			
	id	radius_mean	texture_mean	perimeter_mean	area_mean	\				<class 'pandas.core.frame.DataFrame'>			
count	5.690000e+02	569.000000	569.000000	569.000000	569.000000	\				RangeIndex: 569 entries, 0 to 568			
mean	3.037183e+07	14.127292	19.289649	91.969033	654.889104	\				Data columns (total 33 columns):			
std	1.250206e+08	3.524049	4.301036	24.298981	351.914129	\				#	Column	Non-Null Count	Dtype
min	8.670000e+03	6.981000	9.710000	43.790000	143.500000	\				---	-----	-----	----
25%	8.692180e+05	11.700000	16.170000	75.170000	420.300000	\				0	id	569 non-null	int64
50%	9.060240e+05	13.370000	18.840000	86.240000	551.100000	\				1	diagnosis	569 non-null	object
75%	8.813129e+06	15.780000	21.800000	104.100000	782.700000	\				2	radius_mean	569 non-null	float64
max	9.113205e+08	28.110000	39.280000	188.500000	2501.000000	\				3	texture_mean	569 non-null	float64
	smoothness_mean	compactness_mean	concavity_mean	concave	points_mean	\				4	perimeter_mean	569 non-null	float64
count	569.000000	569.000000	569.000000	569.000000	569.000000	\				5	area_mean	569 non-null	float64
mean	0.096360	0.104341	0.088799	0.048919	0.048919	\				6	smoothness_mean	569 non-null	float64
std	0.014064	0.052813	0.079720	0.038803	0.038803	\				7	compactness_mean	569 non-null	float64
min	0.052630	0.019380	0.000000	0.000000	0.000000	\				8	concavity_mean	569 non-null	float64
25%	0.086370	0.064920	0.029560	0.020310	0.020310	\				9	concave points_mean	569 non-null	float64
50%	0.095870	0.092630	0.061540	0.033500	0.033500	\				10	symmetry_mean	569 non-null	float64
75%	0.105300	0.130400	0.130700	0.074000	0.074000	\				11	fractal_dimension_mean	569 non-null	float64
max	0.163400	0.345400	0.426800	0.201200	0.201200	\				12	radius_se	569 non-null	float64
	symmetry_mean	...	texture_worst	perimeter_worst	area_worst	\				13	texture_se	569 non-null	float64
count	569.000000	...	569.000000	569.000000	569.000000	\				14	perimeter_se	569 non-null	float64
mean	0.181162	...	25.677223	107.261213	880.583128	\				15	area_se	569 non-null	float64
std	0.027414	...	6.146258	33.602542	569.356993	\				16	smoothness_se	569 non-null	float64
min	0.106000	...	12.020000	50.410000	185.200000	\				17	compactness_se	569 non-null	float64
25%	0.161900	...	21.080000	84.110000	515.300000	\				18	concavity_se	569 non-null	float64
50%	0.179200	...	25.410000	97.660000	686.500000	\				19	concave points_se	569 non-null	float64
75%	0.195700	...	29.720000	125.400000	1084.000000	\				20	symmetry_se	569 non-null	float64
max	0.304000	...	49.540000	251.200000	4254.000000	\				21	fractal_dimension_se	569 non-null	float64
	smoothness_worst	compactness_worst	concavity_worst	\				22	radius_worst	569 non-null	float64
count	569.000000	569.000000	569.000000	569.000000	569.000000	\				23	texture_worst	569 non-null	float64
mean	0.132369	0.254265	0.272188	0.272188	0.272188	\				24	perimeter_worst	569 non-null	float64
std	0.022832	0.157336	0.208624	0.208624	0.208624	\				25	area_worst	569 non-null	float64
min	0.071170	0.027290	0.000000	0.000000	0.000000	\				26	smoothness_worst	569 non-null	float64
25%	0.116600	0.147200	0.114500	0.114500	0.114500	\				27	compactness_worst	569 non-null	float64
50%	0.131300	0.211900	0.226700	0.226700	0.226700	\				28	concavity_worst	569 non-null	float64
75%	0.146000	0.339100	0.382900	0.382900	0.382900	\				29	concave points_worst	569 non-null	float64
max	0.222600	1.058000	1.252000	1.252000	1.252000	\				30	symmetry_worst	569 non-null	float64
	concave points_worst	symmetry_worst	fractal_dimension_worst	\				31	fractal_dimension_worst	569 non-null	float64
count	569.000000	569.000000	569.000000	569.000000	569.000000	\				32	Unnamed: 32	0 non-null	float64
mean	0.114606	0.290076	0.083946	0.083946	0.083946	\				dtypes: float64(31), int64(1), object(1)			
std	0.065732	0.061867	0.018061	0.018061	0.018061	\				memory usage: 146.8+ KB			
min	0.000000	0.156500	0.055040	0.055040	0.055040	\							
25%	0.064930	0.250400	0.071460	0.071460	0.071460	\							
50%	0.099930	0.282200	0.080040	0.080040	0.080040	\							
75%	0.161400	0.317900	0.092080	0.092080	0.092080	\							
max	0.291000	0.663800	0.207500	0.207500	0.207500	\							

after that below phases has been followed.

Phase 1: Exploratory Data Analysis (EDA) & Preprocessing

Exploratory Data Analysis (EDA) helps understand the dataset, detect patterns, and identify potential issues before model training. Below is a step-by-step guide to performing EDA on the selected.

- Check missing values and handle them (Found none, except an extra column, which was dropped)

```
#Checking null values
print(data.isnull().sum)
```

```
<bound method DataFrame.sum of
```

	id	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	\
0	False	False	False	False	False	False	
1	False	False	False	False	False	False	
2	False	False	False	False	False	False	
3	False	False	False	False	False	False	
4	False	False	False	False	False	False	
..	
564	False	False	False	False	False	False	
565	False	False	False	False	False	False	
566	False	False	False	False	False	False	
567	False	False	False	False	False	False	
568	False	False	False	False	False	False	


	smoothness_mean	compactness_mean	concavity_mean	concave points_mean	\
0	False	False	False	False	
1	False	False	False	False	
2	False	False	False	False	
3	False	False	False	False	
4	False	False	False	False	
..	
564	False	False	False	False	
565	False	False	False	False	
566	False	False	False	False	
567	False	False	False	False	
568	False	False	False	False	

	radius_worst	texture_worst	perimeter_worst	area_worst	\
0	...	False	False	False	
1	...	False	False	False	
2	...	False	False	False	

➤ Visualize class distribution

Visualize class distribution

```
[ ] #Plotting countplot for visualizing the class distribution
sns.countplot(x='diagnosis', data=data)
plt.title("class Distribution")
plt.show()
```




➤ Perform correlation analysis (Correlation has been plotted, while fixing mix-up of values)

Perform correlation analysis

```
[ ] #For Label encoding our target to analyze its corr
label_encoder = LabelEncoder()

#Fitting encoder to diagnosis column for transforming it as it is the target
data['diagnosis_encoded'] = label_encoder.fit_transform(data['diagnosis'])

[ ] #Corr Analysis
plt.figure(figsize=(15, 10))
numeric_data = data.select_dtypes(include=np.number)
sns.heatmap(numeric_data.drop(columns=['id']).corr(), annot=True, cmap='coolwarm', annot_kws={'size': 8})
plt.title("Correlation Analysis")
plt.xticks(rotation=45, ha='right') # Rotate x-axis labels
plt.tight_layout() # Adjust layout to prevent overlapping labels
plt.show()
```








- **Identify outliers** (Outliers been identified via box-plot, some of the plots are pasted below, rest could be checked from the notebook submitted)

▼ Identify outliers

```
#Identifying outliers
numerical_features = data.select_dtypes(include=np.number).drop(columns=['id'])

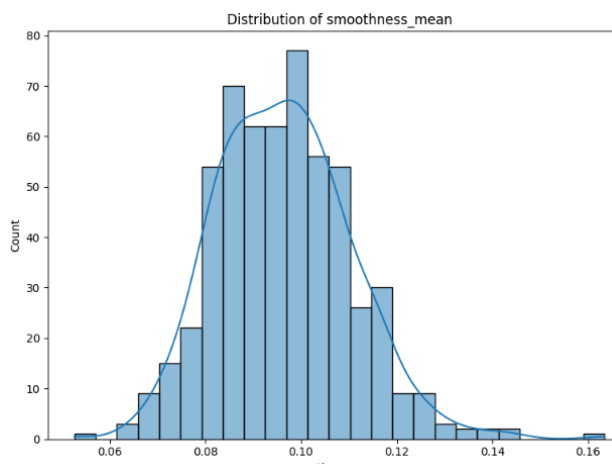
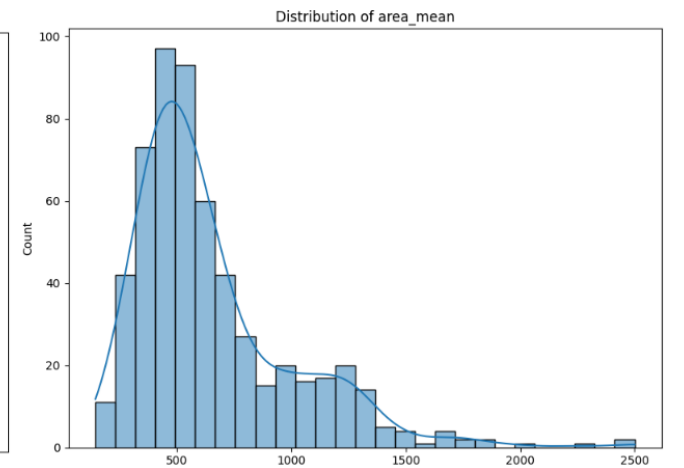
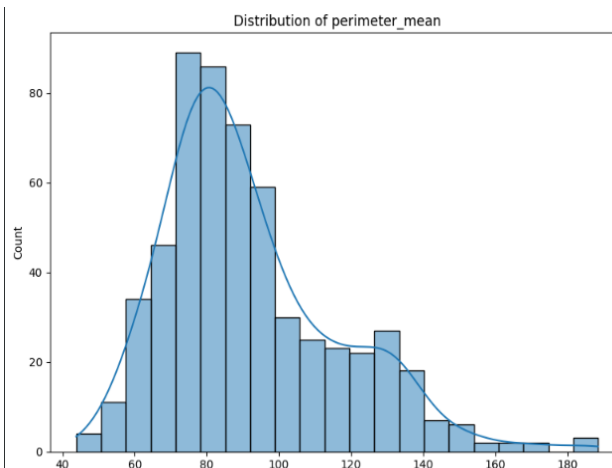


#Creating boxplots for each feature individually
for feature in numerical_features.columns:
    plt.figure(figsize=(10, 8))
    data.boxplot(column=feature)
    plt.title(f'Boxplot for {feature}')
    plt.xticks(rotation=45, ha='right')
    plt.tight_layout()
    plt.show()
```



- **Visualize feature distributions** (Feature distribution has been visualised; some of the distributions are here under; the rest could be analysed from the notebook)

Visualize feature distributions

```
[ ] #Visualizing feature distributions
for feature in numerical_features.columns:
    plt.figure(figsize=(8, 6))
    sns.histplot(data[data[feature]], kde=True)
    plt.title(f'Distribution of {feature}')
    plt.tight_layout();
```



- **Normalize the dataset** (The dataset has been normalized by via standard scaler using fit transform)

```

  ▾ Normalize the dataset

[ ] #selectign numerical columns for scalling
    numerical_features = data.select_dtypes(include=np.number).drop(columns=['id'])

    #Normalizing the dataset
    scaler = StandardScaler()
    data[numerical_features.columns] = scaler.fit_transform(data[numerical_features.columns])

```

- **Split data into training & testing sets** (The data has been split into training and testing sets)

```

  ▾ Split data into training & testing sets

[ ] #Splitting data into training and testing sets
    X = data.drop(columns=['id', 'diagnosis', 'diagnosis_encoded'])
    y = data['diagnosis_encoded']

[ ] # Normalize features
    scaler = StandardScaler()
    X_scaled = scaler.fit_transform(X)

[ ] # Split data into training and test sets
    X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, random_state=42)

```

Phase 2: PCA for dimensionality reduction

- **Analyze the impact of reducing the number of features on classification performance.**
- **Compute the explained variance ratio for different numbers of principal components.**
- **Select the number of components that retain 95% of variance.**

The PCA has been applied for dimensionality reduction, the variance ratio has been depicted, and components containing 95pc of variance have been selected.

```

  ▾ Phase 2: PCA for dimensionality reduction

Analyze the impact of reducing the number of features on classification performance.

Compute the explained variance ratio for different number of principal components.

  ▾ Select the number of components that retain 95% of variance.

[ ] #Analyzing impact of reducing no. of features
    from sklearn.decomposition import PCA

    pca = PCA(n_components=0.95) # Retain 95% of the variance
    X_train_pca = pca.fit_transform(X_train)
    X_test_pca = pca.transform(X_test)

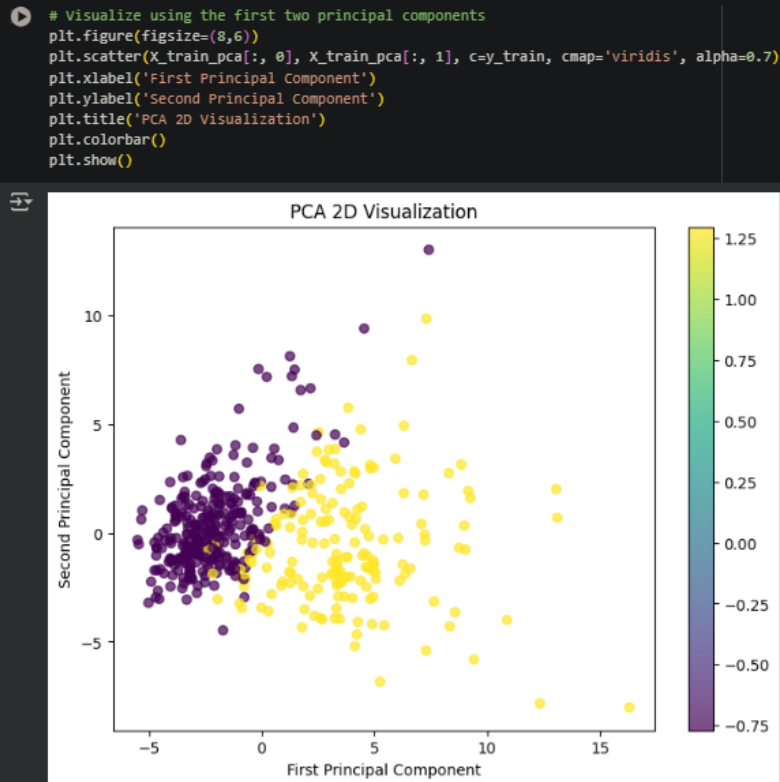
    # Explained variance ratio
    print("Explained variance ratio:", pca.explained_variance_ratio_)
    print("Number of components selected:", pca.n_components_)

Explained variance ratio: [0.4325482  0.19732988 0.09865421 0.0625777  0.053537   0.04102624
 0.02274171 0.01636683 0.01404727 0.01198947]
Number of components selected: 10

```


- **Visualize the data in 2D using the first two principal components.** (The visualization of the data has been done using first two principal components)

Visualize the data in 2D using the first two principal components.



Comparison:

The comparison of the models has been carried out, and the results are as under:

- **SVM**

Support Vector Machine (SVM)

```
from sklearn.svm import SVC
from sklearn.metrics import classification_report, confusion_matrix, roc_auc_score

#Converting to integers representing classes
y_train = y_train.astype(int)
y_test = y_test.astype(int)

# List of hyperparameters to try
kernels = ['linear', 'rbf', 'poly', 'sigmoid']
C_values = [0.01, 0.1, 1, 10, 100]
gammas = ['scale', 'auto', 0.001, 0.01, 0.1, 1]

best_score = 0
best_params = {}

# Grid search over a subset of hyperparameters
for kernel in kernels:
    for C in C_values:
        svm_model = SVC(kernel=kernel, C=C, gamma='scale', probability=True, random_state=42)
        svm_model.fit(X_train, y_train) # Now using categorical target
        score = svm_model.score(X_test, y_test)
        if score > best_score:
            best_score = score
            best_params = {'kernel': kernel, 'C': C}

print("Best SVM parameters (original):", best_params)
print("Best SVM accuracy (original):", best_score)

# Evaluating the best SVM on PCA-transformed data
svm_pca = SVC(kernel=best_params['kernel'], C=best_params['C'], gamma='scale', probability=True, random_state=42)
svm_pca.fit(X_train_pca, y_train)
score_pca = svm_pca.score(X_test_pca, y_test)
print("Best SVM accuracy (PCA-transformed):", score_pca)
```

Best SVM parameters (original): {'kernel': 'linear', 'C': 0.1}
Best SVM accuracy (original): 0.9824561403508771
Best SVM accuracy (PCA-transformed): 0.9824561403508771

```

Evaluation Metrics

[ ] y_pred = svm_model.predict(X_test)
    print(classification_report(y_test, y_pred))
    print(confusion_matrix(y_test, y_pred))

precision    recall  f1-score   support

     0       0.96       0.97       0.97        71
     1       0.95       0.93       0.94         43

 accuracy          0.96          0.96          0.96          114
 macro avg          0.96          0.95          0.95          114
 weighted avg          0.96          0.96          0.96          114

[[69  2]
 [ 3 40]]

```

➤ Decision Tree

```

Decision Tree

from sklearn.tree import DecisionTreeClassifier, export_graphviz
import graphviz

# Without pruning
dt_model = DecisionTreeClassifier(criterion='gini', random_state=42)
dt_model.fit(X_train, y_train)
print("Decision Tree (unpruned) accuracy:", dt_model.score(X_test, y_test))

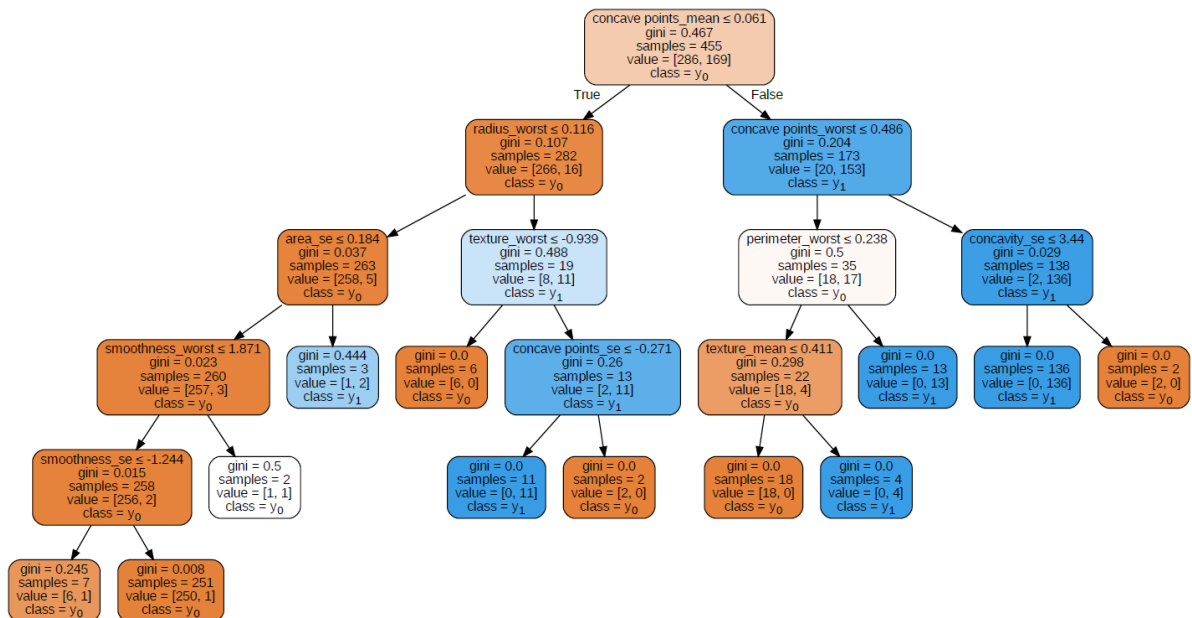
# With pruning parameters (pre-pruning)
dt_pruned = DecisionTreeClassifier(criterion='gini', max_depth=5, min_samples_split=5, min_samples_leaf=2, random_state=42)
dt_pruned.fit(X_train, y_train)
print("Decision Tree (pruned) accuracy:", dt_pruned.score(X_test, y_test))

# Visualize the pruned tree
dot_data = export_graphviz(dt_pruned, out_file=None,
                           feature_names=X.columns,
                           class_names=True,
                           filled=True, rounded=True,
                           special_characters=True)

graph = graphviz.Source(dot_data)
graph.render("decision_tree")

Decision Tree (unpruned) accuracy: 0.9473684210526315
Decision Tree (pruned) accuracy: 0.956140350877193
'decision_tree.pdf'

```




```
▼ Evaluation

y_pred_dt = dt_pruned.predict(X_test)
print(classification_report(y_test, y_pred_dt))
print(confusion_matrix(y_test, y_pred_dt))
```

		precision	recall	f1-score	support
	0	0.96	0.97	0.97	71
	1	0.95	0.93	0.94	43
	accuracy			0.96	114
	macro avg	0.96	0.95	0.95	114
	weighted avg	0.96	0.96	0.96	114

```
[[69  2]
 [ 3 40]]
```

➤ Neural Network

```
▼ Neural Network

import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout
from tensorflow.keras.optimizers import Adam

# Define the model architecture (using original data)
model = Sequential([
    Dense(64, activation='relu', input_shape=(X_train.shape[1],)),
    Dropout(0.2),
    Dense(32, activation='relu'),
    Dense(1, activation='sigmoid') # For binary classification
])

model.compile(optimizer=Adam(learning_rate=0.001), loss='binary_crossentropy', metrics=['accuracy'])
history = model.fit(X_train, y_train, epochs=100, batch_size=32, validation_split=0.2, verbose=0)
loss, accuracy = model.evaluate(X_test, y_test)
print("Neural Network (original) accuracy:", accuracy)

# Using PCA-transformed data
model_pca = Sequential([
    Dense(64, activation='relu', input_shape=(X_train_pca.shape[1],)),
    Dropout(0.2),
    Dense(32, activation='relu'),
    Dense(1, activation='sigmoid')
])

model_pca.compile(optimizer=Adam(learning_rate=0.001), loss='binary_crossentropy', metrics=['accuracy'])
history_pca = model_pca.fit(X_train_pca, y_train, epochs=100, batch_size=32, validation_split=0.2, verbose=0)
loss_pca, accuracy_pca = model_pca.evaluate(X_test_pca, y_test)
print("Neural Network (PCA-transformed) accuracy:", accuracy_pca)
```

```
/usr/local/lib/python3.11/dist-packages/keras/src/layers/core/dense.py:87: UserWarning: Do not pass an `input_shape`/'input_dim` argument to a layer.
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
4/4 [100%] 0s 15ms/step - accuracy: 0.9651 - loss: 0.1999
Neural Network (original) accuracy: 0.9649122953414917
4/4 [100%] 0s 9ms/step - accuracy: 0.9558 - loss: 0.1994
Neural Network (PCA-transformed) accuracy: 0.9649122953414917
```

Phase 3: Model Comparison & Final Report

➤ Model Comparison

```
▼ Model Comparison

[ ] # Model Comparison
models = {
    "SVM": svm_model,
    "SVM (PCA)": svm_pca,
    "Decision Tree": dt_pruned,
    "Neural Network": model,
    "Neural Network (PCA)": model_pca
}

results = []
for name, model in models.items():
    if name.endswith("(PCA)"):
        y_pred = model.predict(X_test_pca)
    else:
        y_pred = model.predict(X_test)

    if name in ["Neural Network", "Neural Network (PCA)":
        y_pred_binary = (y_pred > 0.5).astype(int) # Convert probabilities to binary predictions
    else:
        y_pred_binary = y_pred

    report = classification_report(y_test, y_pred_binary, output_dict=True)
    results.append({
        "Model": name,
        "Accuracy": report['accuracy'],
        "Precision": report['weighted avg']['precision'],
        "Recall": report['weighted avg']['recall'],
        "F1-score": report['weighted avg']['f1-score'],
        "ROC-AUC": roc_auc_score(y_test, y_pred) if name in ["SVM", "SVM (PCA)", "Neural Network", "Neural Network (PCA)"] else roc_auc_score(y_test, y_pred_binary) # Calculate ROC-AUC if applicable
    })

results_df = pd.DataFrame(results)
print(results_df)
```

The SVM and Neural Network has been the top performers among all.

4/4		0s 26ms/step				
4/4		0s 29ms/step				
	Model	Accuracy	Precision	Recall	F1-score	ROC-AUC
0	SVM	0.956140	0.956088	0.956140	0.956036	0.951032
1	SVM (PCA)	0.982456	0.982456	0.982456	0.982456	0.981330
2	Decision Tree	0.956140	0.956088	0.956140	0.956036	0.951032
3	Neural Network	0.964912	0.964912	0.964912	0.964912	0.992794
4	Neural Network (PCA)	0.964912	0.965828	0.964912	0.965060	0.992466

➤ Confusion Metrices

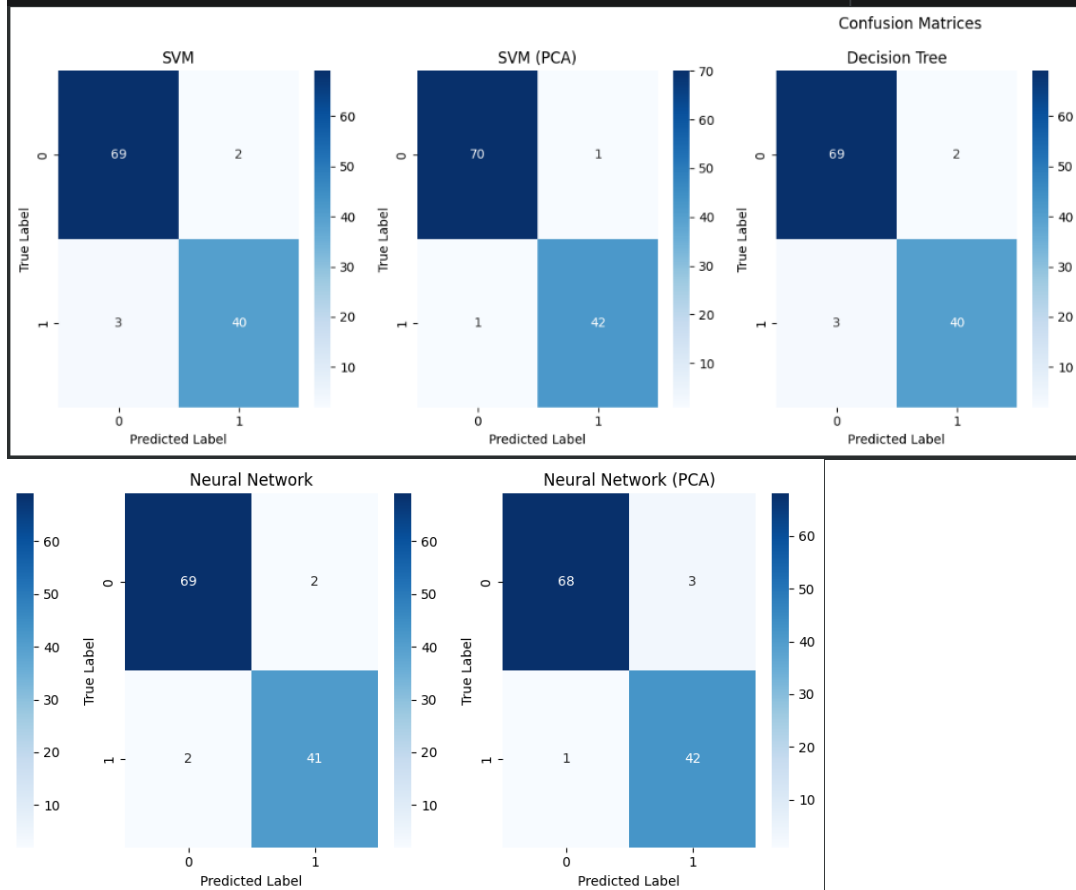
```
[ ] # Confusion Matrices
fig, axes = plt.subplots(1, len(models), figsize=(20, 5))
fig.suptitle("Confusion Matrices")

for i, (name, model) in enumerate(models.items()):
    if name.endswith("(PCA)":
        y_pred = model.predict(X_test_pca)
    else:
        y_pred = model.predict(X_test)

    if name in ["Neural Network", "Neural Network (PCA)":
        y_pred_binary = (y_pred > 0.5).astype(int) # Convert probabilities to binary predictions
    else:
        y_pred_binary = y_pred

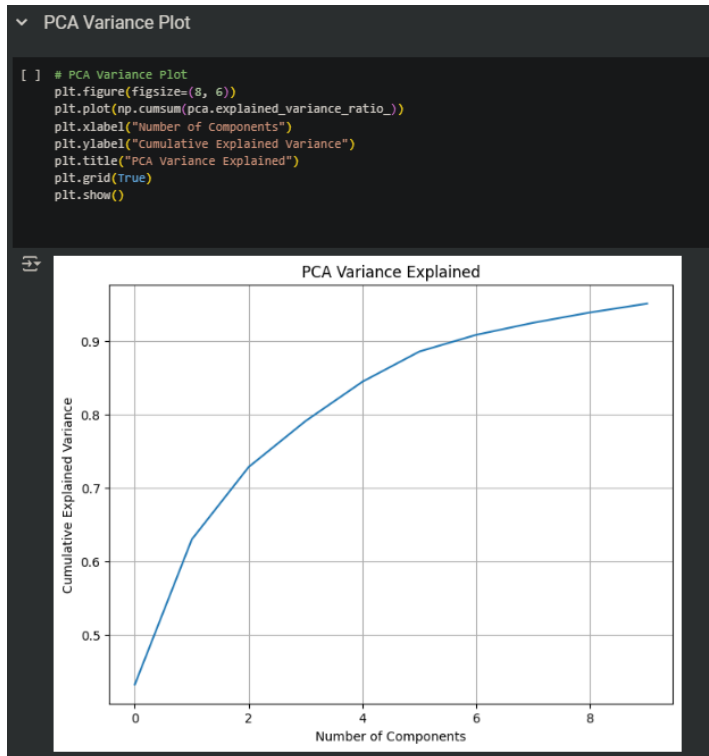
    cm = confusion_matrix(y_test, y_pred_binary)
    sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", ax=axes[i])
    axes[i].set_title(name)
    axes[i].set_xlabel("Predicted Label")
    axes[i].set_ylabel("True Label")

plt.tight_layout()
plt.show()
```

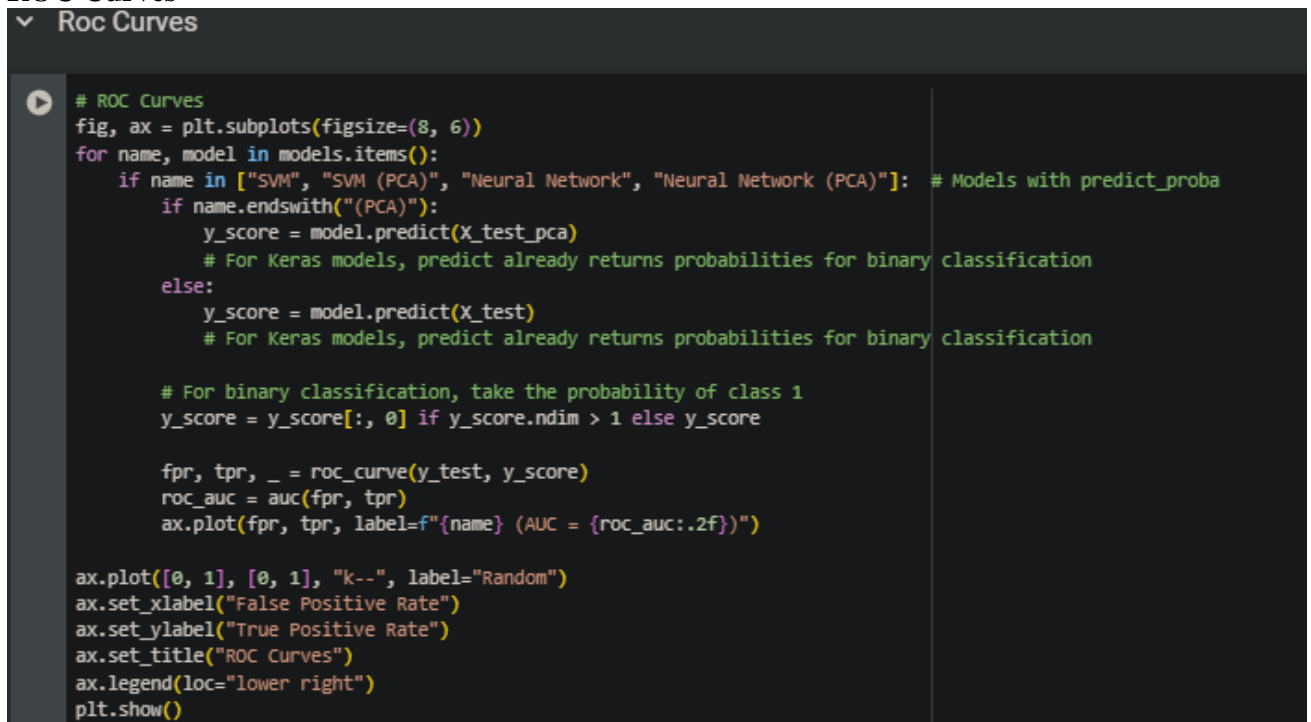


➤ PCA Variance Plot


After analyzing the plot it can be clearly seen that first few principal components explain significant portion of variance suggesting that PCA is likely to be effective without losing much information.



➤ ROC Curves



The ROC curves have been plotted with their corresponding Area Under the Curve (AUC) values. The SVM and Neural Network models appear to be the top performers, which has already been conveyed while model comparisons.



Conclusion:

1. Which model performed best?

The SVM and Neural Network models (both with and without PCA) achieved the best results in terms of accuracy, precision, recall, and F1-score. SVM with PCA showed similar performance to SVM without PCA, though its scores were slightly lower. The Neural Network with PCA did not significantly impact performance and performed comparably to the original model without PCA.

2. How did PCA impact classification performance?

In this case, PCA had minimal effect on classification performance and did not lead to significant improvements. While PCA can enhance model performance in certain situations, it seems that the original features contained crucial information that was not compromised during dimensionality reduction with PCA.

3. Which approach is most suitable for medical diagnosis?

Although SVM and Neural Networks (both with and without PCA) performed similarly, Neural Networks might be slightly more suitable for this specific case due to their higher accuracy and other metrics, especially when compared to SVM without PCA. However, all the models performed comparably, allowing us flexibility in our choice.